

Benchmark Radar v0.9.0

From daily collection to benchmark search and score history

Koutian Wu

29 August 2026 | Software v0.9.0 | Data cutoff 2026-08-29 | Git 98c7de3

Reserved DOI: 10.5281/zenodo.22167102

7,540 source observations across 37 snapshots	4,537 unique artifacts in the cumulative evidence graph	1,242 searchable entries across 4 search sources	37 public collection sources monitored
--	--	---	---

Benchmark Radar brings daily discoveries, catalog search, vendor reporting, and sourced model scores into one place. Each layer keeps its own count because each records a different kind of evidence.

Executive findings

- Search covers 1,242 entries. Three external catalogs supply 1,173 rows. The curated score archive adds 69 benchmark tracks.
- The daily corpus holds 7,540 observations for 4,537 artifacts. Primary or structured sources supplied 7,523 observations, and each record keeps its source URL and retrieval metadata.
- Frozen v0.9.0 population: 36 model reports cover 94 curated benchmarks. Those cutoff counts remain the deposited report baseline. The issue #456 addendum separately audits 37 reviewed documents, 12 organization labels, and 110 benchmark IDs pinned at commit 98c8cf6.
- Identity and task classification need more work. Forty-one artifacts link across multiple sources. KW-Bench has produced 4,129 tracks, with classification still empty in v0.9.0.

Abstract

Benchmark Radar collects benchmark work from 37 public sources and publishes it through a dashboard, RSS, JSON, and an offline client. This report checks the full path from collection to search. It covers source health, exact-identifier matching, the ranking rubric, the 1,242-entry search index, model-card adoption, score histories, and the reports already in the repository. The system has strong provenance and reproducible builds. Cross-source matching, protocol capture, semantic search, and KW-Bench classification remain active work.

1. Product tour

The README leads with the user task: find a benchmark in seconds, then inspect model-card adoption and score movement. The dashboard and RSS feed serve readers. JSON, the local CLI, and the local HTTP API support reproducible research.

Benchmark Radar's daily-intelligence logic was inspired by BuilderPulse, an AI-powered daily intelligence project for indie hackers and builders that answers 20 questions from 10+ sources every morning [8]. Benchmark Radar adapts that multi-source briefing pattern to benchmark discovery, normalized catalog search, and score history.

Search every benchmark

Search benchmarks, tasks, domains...

1,242 benchmarks · 4 sources

1.1 Search coverage

Search layer	Rows	Best field	Score data
LLM Stats	687	Benchmark names and descriptions	5,544 rows; evaluation setup is sparse
OpenCompass Hub	461	Paper, repository, release, and dataset links	Historical leaderboards stay source-labelled
Artificial Analysis	25	Current commercial evaluation catalog	7,050 third-party rows
Curated score tracks	69	Scores read from primary model reports	285 rows with instrument and protocol fields
Search box	1,242	All four layers in one index	Source label shown on every result

The UI adds 1,173 external rows and 69 curated score tracks. It keeps source records separate, so the same benchmark name can appear in more than one catalog. The label “4 sources” refers to these four search layers. The daily collector uses 37 public endpoints.

1.2 Dashboard and local tools

Surface	Reader question	Evaluation
Today	What appeared or changed?	Ranks new records and links the source evidence.
Search	Which records match this name, task, or domain?	Ranks exact names, phrases, and field-token matches.
Leaderboard	Which benchmarks do labs report?	Counts benchmark mentions across vendor reports.
Scores	Which printed values can be connected?	Matching instrument + protocol creates a series.
Trends / map	Which topics and sources recur?	Coverage signatures gate comparisons.
CLI + HTTP	Can an analyst reproduce search offline?	One QueryService and stable JSON contract.

2. Pipeline evaluation



The same snapshots and code produce the same derived files.

2.1 Collection and health

Each run queries enabled connectors inside a 48-hour window, records counts and errors, drops future-dated rows, then requires the configured core sources to be healthy. arXiv, Hugging Face, and GitHub are required. On the cutoff run all three were healthy; Semantic Scholar returned HTTP 429 and Brave had no API key; OpenReview was healthy but empty.

The cutoff run fetched 826 rows, deduplicated them to 783, classified 273 as eligible, and recommended 97. Two collections ran that day. Their merged page contains 528 records, including 186 recommendations.

2.2 Normalization, identity, and retention

- Stable fields include source ID, URL, title, timestamps, authors or organizations when supplied, parser version, retrieval time, and a SHA-256 payload fingerprint. The publisher omits raw responses and credentials.
- The resolver merges observations that share a DOI, arXiv ID, OpenReview ID, GitHub repository, or Hugging Face artifact. A similar title leaves two records in place until a reviewer confirms the match.
- The corpus includes eligible records below the recommendation threshold. The recommended tag controls display priority.

2.3 Interpretation, publication, and query

Priority combines relevance (35%), evidence (20%), recency (20%), and adoption (25%) on a 0–100 scale. Scoring version 5 ships with the data. Generators replay validated snapshots into the cumulative graph, normalize external catalogs, classify KW-Bench tracks, and package a checksummed release. Installed clients switch to a new data version after checksum and schema validation. Search reads the active local version.


Area	Working today	Current gap
Reproducibility	Snapshots, schemas, parser versions, hashes, deterministic rebuilds.	Four early snapshots are simulated.
Reliability	Required-source gate; optional failures visible; atomic client activation.	Rate limits and secrets change realized coverage.
Identity	Exact-anchor merges and reviewed external groups.	Forty-one artifacts have more than one source.
Interfaces	CLI and HTTP share QueryService and JSON.	Local token matching misses some paraphrases.
Verification	Clean-worktree rebuild plus a passing full CI suite.	Source coverage still depends on public endpoints.

3. Counts and units

The dashboard publishes several counts because it tracks several units. Use the unit and cutoff beside the number.

Unit / product	Count	Definition	Use it for
Snapshot	37	33 observed, 4 simulated.	The dated history available for replay.
Source observation	7,540	Persisted source record.	Evidence volume in cumulative replay.
Unique artifact	4,537	Exact-ID-resolved paper, repo, dataset, release, or page.	Distinct artifacts observed.
External catalog row	1,173	One row from three catalog sources.	Broad, source-labelled catalog search.
Searchable entry	1,242	1,173 external + 69 curated score tracks.	Current web-search reach.
Adoption benchmark (frozen v0.9.0)	94	Canonical identity in curated registry.	Model-card adoption and missing adoption.
Model-card document (frozen v0.9.0)	36	Curated reports from 11 organizations.	Documents read for mentions.
Curated score	285	Numeric result from a cited document; 69 tracks.	Score histories grouped by evaluation setup.
Model registry	861	Models across curated and crawled layers; 19 in both.	Finding a model across both data layers.

3.1 Cumulative source composition



7,540 cumulative source observations through 29 August 2026

Five normalized source labels supply 7,376 of 7,540 observations (97.8%). The other endpoints add discovery routes and publish their own yield and health status.

3.2 Adoption and score findings

The issue #456 sensitivity audit rebuilds vendor-reporting counts from the reviewed registry rather than repeating the frozen PDF prose. Its primary full-history rule finds 16 canonical benchmark IDs reported by at least six organizations, while the pre-registered document, time-window, family, and missing-report alternatives produce different boundaries. Section 6.1 reports the result and limitations. The score archive remains separate: it finds eight bounded metrics with five points of headroom or less, and six benchmarks gained model-card mentions after their last readable score.

4. Source inventory and health

The daily collector reads 12 discovery connectors, 24 first-party feeds, and Hacker News. The search box draws from four catalog layers.

4.1 Direct discovery connectors (12)

Connector	Role	Cutoff run	Core?
arXiv	Primary papers via cs.AI/CL/CV/SE RSS	Healthy · 23	Yes
Hugging Face Hub	Datasets and Spaces	Healthy · 126	Yes
GitHub Search	Code and artifacts	Healthy · 300; at cap	Yes
GitHub Organizations	Reviewed organization repos	Healthy · 15	No
Hugging Face Papers	Community-surfaced papers	Healthy · 23	No
Kaggle Datasets	Public benchmark datasets	Healthy · 28	No
Zenodo	DOI-bearing artifacts	Healthy · 77	No
OpenReview	Conference submissions	Healthy · 0	No
Semantic Scholar	Structured scholarly discovery	Failed · HTTP 429	No
GitHub Releases	Curated first-party releases	Healthy · 3	No
OpenAlex	Scholarly discovery	Healthy · 228	No
Brave Search	Web and official domains	Unavailable · no key	No

4.2 First-party feeds (24)

Feeds 1–12	Feeds 13–24
Meituan Engineering OpenAI News Google AI Google DeepMind Google Research Apple Machine Learning Research AWS Machine Learning Hugging Face Blog Microsoft Research NVIDIA AI Blog Mistral AI Meta Research	Ai2 Together AI Sakana AI Qwen Ollama Stability AI Nomic AI Replicate NVIDIA Developer IBM Research Databricks LangChain

The feed collector yielded three relevant records. Some healthy feeds produced zero matches after relevance filtering. Domain-constrained searches cover organizations that lack a verified feed. The collector excludes third-party mirrors.

4.3 Public attention (1)

The Hacker News collector uses the public Algolia API. It found 12 records on the cutoff run and places them in a separate attention feed. The priority calculation uses evidence sources only.

5. Data quality and report history

Finding	Evidence	Interpretation
High provenance	7,523 / 7,540 primary or structured (99.77%).	Readers can open the upstream record for almost every observation.
Low corroboration	41 / 4,537 artifacts have >1 normalized source (0.90%).	Exact identity prevents bad joins but leaves plausible duplicates.
Uneven yield	Five source labels provide 97.8% of observations.	Caps or outages can change apparent topic mix.
Simulation	23–26 July are simulated snapshots.	Use them to test replay and UI history.
Classification gap	KW-Bench: 4,129 tracks, 0 classified.	The task-capability view has no classifications yet.
Protocol sparsity	12,594 aggregator scores stay outside curated progression.	Comparison needs shots, harness, tools, attempts, and date.
Lexical retrieval	Field tokens drive rank; no semantic reranker.	Reproducible, but paraphrases can be missed.

5.1 Existing reports

Document	Use today	Date limit
Landscape report, 31 Jul	A dated study of 791 sightings, 645 artifacts, and 78 agentic-evaluation artifacts.	Its totals stop on 31 July, before the external catalog and newer connectors.
Source probes, 27 Aug	Verified HF Papers, Kaggle, Spaces, and Zenodo endpoints.	The probes record one check on 27 August.
External catalog audit	OpenCompass identity fields and the score-heavy peer catalogs.	The audit keeps same-name records separate until identity review.
Daily report / briefing	Triage with citations and source health.	Each edition covers one collection window.

5.2 Reading the evidence

- You can trace a new or updated record to its source, inspect why the taxonomy matched it, and find the vendor document behind each curated benchmark mention.
- For score comparisons, match the instrument and protocol fields. For trend comparisons, use days with the same connector coverage and report limit. Market-size estimates need a separate benchmark-family census.

6. Findings from the current data

Benchmark Radar turns papers, repositories, datasets, and model reports into records you can search, filter, download, and query offline. Three results stand out in the current data.

6.1 A recurring reporting group with a definition-sensitive boundary

Definition	Observed threshold set
Full reviewed history	16 canonical IDs · 37 documents · 12 organizations
Trailing 365 days	6 IDs · same 6-organization threshold
Latest document per organization	3 IDs · 12 documents
Reviewed family projection	13 families · score tracks unmerged

The original eight-item statement does not reproduce from its stated rule. Across all 37 reviewed documents from 12 organization labels, 16 canonical benchmark IDs—not eight—appear in documents from at least 6 organizations. The eight names printed in the prior draft are a strict subset of that threshold set and match a ranking truncation, not a complete threshold-defined core.

The boundary changes under the pre-registered alternatives. The trailing 365-day window contains 6 IDs at the same threshold; one latest document per organization contains 3; model-card documents alone contain 4; and the explicit reviewed-family projection contains 13 families. These results support a narrower claim that a recurring reporting group exists in the reviewed sample. They do not support an exact eight-benchmark boundary or a field-wide statement that vendor attention has converged.

Method and limits: Junkai Wang / @JunkaiWang-TheoPhy contributed the issue #456 audit (<https://github.com/ktwu01/benchmark-radar/issues/456>) at source commit 98c8cf6fb5d1d69c66d438ea9f92242b2205c9ae and cutoff 2026-08-31. Counts are rebuilt from data/model_cards.yml as binary organization-by-benchmark mention edges; documents, models, benchmark IDs, families, and score tracks remain separate. The sensitivity grid varies thresholds, document types, latest-report selection, time windows, reviewed families, and missing-report stress tests. The registry is a reviewed convenience sample rather than a census: a not-observed cell can mean an unread or missing report and is not evidence of vendor omission. Every aggregate links back to model-card IDs and URLs in the issue #456 machine-readable audit tables [9-10].

6.2 Several bounded metrics are near their ceiling

The curated layer records five points of headroom or less for AIME, Arena-Hard, DeepSearchQA, HMMT, MATH-500, MathVision, SWE-bench Verified, and tau2-bench. Read each value with its reasoning budget, tools, attempts, and evaluator. Those settings often explain score movement between model reports.

6.3 Broad search, deeper curation

The external catalog holds 1,173 rows, more than twelve times the 94-benchmark adoption registry. Use catalog search to find candidates. The curated registry adds the model reports, organizations, instruments, and protocols needed for comparison.

6.4 What to build next

Priority	Measurement work	Result
1	Show the count unit beside every UI total and export.	Readers can see which population each number counts.
2	Resolve high-value cross-source identities with anchors and review.	More records gain paper, code, and dataset links.
3	Complete KW-Bench extraction and reviewed coverage.	Turns 4,129 unclassified tracks into a task map.
4	Expand protocol capture from newer model cards.	Closes the mention-versus-score recency gap.
5	Add optional semantic retrieval while retaining lexical reasons.	Finds paraphrases while retaining visible match reasons.

Use it

Cite Benchmark Radar for source-linked benchmark discovery, catalog search, and vendor-reporting evidence. Use the source and protocol fields when you compare records or scores.

7. Reproducibility, access, and citation

This report evaluates Benchmark Radar v0.9.0 at Git commit 98c7de3 and data cutoff 2026-08-29. The clean worktree ran the CI sequence: lint and formatting checks, external normalization, KW-Bench classification, checksummed data-release construction, and the full test suite. The current full CI suite passed.

Artifact	Canonical or permanent location
Technical report	https://doi.org/10.5281/zenodo.22167102
Source code	https://github.com/ktwu01/benchmark-radar
Dashboard	https://benchmark-radar.org/
Cumulative JSON	https://benchmark-radar.org/data/radar.json
Benchmark catalog	https://benchmark-radar.org/data/benchmark-index.json
RSS	https://benchmark-radar.org/feed.xml
Citation metadata	https://github.com/ktwu01/benchmark-radar/blob/main/CITATION.cff

Suggested citation

Wu, K. (2026). Benchmark Radar v0.9.0: Technical Report. Zenodo. <https://doi.org/10.5281/zenodo.22167102>

Data statement

The v0.9.0 core and frozen section 3 rows use commit 98c7de3 and cutoff 2026-08-29. The issue #456 addendum separately audits 37 documents, 12 organization labels, and 110 benchmark IDs from a SHA-256-pinned registry at commit 98c8cf6 and cutoff 2026-08-31. Do not combine these populations; cite rolling dashboard values with a retrieval date.

References

- [1] K. Wu. Benchmark Radar, version 0.9.0. GitHub, 2026. <https://github.com/ktwu01/benchmark-radar>
- [2] K. Wu. AI Benchmark Landscape Report. 2026. <https://github.com/ktwu01/benchmark-radar/blob/main/docs/reports/ai-benchmark-landscape-report.md>
- [3] K. Wu. Benchmark Radar cumulative corpus schema. 2026. <https://github.com/ktwu01/benchmark-radar/blob/main/docs/cumulative-corpus.schema.json>
- [4] K. Wu. Source probe evidence. 2026. <https://github.com/ktwu01/benchmark-radar/blob/main/docs/source-probe-evidence.md>
- [5] D. S. Katz et al. Recognizing the value of software: a software citation guide. F1000Research 9:1257, 2021. <https://doi.org/10.12688/f1000research.26932.2>
- [6] A. Smith, D. S. Katz, and K. E. Niemeyer. Software citation principles. PeerJ Computer Science 2:e86, 2016. <https://doi.org/10.7717/peerj-cs.86>
- [7] Citation File Format developers. Citation File Format 1.2.0. <https://citation-file-format.github.io/>
- [8] L. Xiaopai. BuilderPulse: AI-powered daily intelligence for indie hackers and builders. GitHub, 2026. <https://github.com/BuilderPulse/BuilderPulse>
- [9] Benchmark Radar contributors. Reviewed model-card benchmark registry. 2026. https://github.com/ktwu01/benchmark-radar/blob/main/data/model_cards.yml
- [10] J. Wang. Vendor-attention sensitivity audit for issue #456. 2026. <https://github.com/ktwu01/benchmark-radar/tree/main/docs/technical-report/vendor-attention-audit>

Software: MIT License.

Technical report and original editorial content: CC BY-NC 4.0.

Commercial republication, resale, paid newsletters, dataset packaging, or commercial product integration requires prior written permission from Koutian Wu.

Third-party source material remains under its original terms.