

Benchmark Radar

An evidence-first index for tracking AI benchmarks

Koutian Wu

29 August 2026 | Software v0.9.0 | Report v1.0

DOI: 10.5281/zenodo.22167102

7,540 source observations across 37 snapshots	4,537 unique artifacts in the cumulative evidence graph	1,173 searchable benchmark catalog records	36 model cards and release documents
--	--	---	---

Benchmark Radar helps an evaluator find a new benchmark, inspect the source that announced it, and check how vendors use it without treating adoption as proof of quality.

Abstract

Researchers now encounter new evaluation suites across papers, code repositories, datasets, model cards, and release posts. Benchmark Radar collects benchmark-related signals from 37 public sources each day, normalizes them into evidence-linked records, and publishes the result through a dashboard, RSS feed, JSON corpus, and offline query client. Through 29 August 2026, the cumulative corpus contains 7,540 source observations covering 4,537 unique artifacts across 37 snapshots; a separate catalog contains 1,173 searchable benchmark records. The project also maintains a curated model-card registry and a protocol-aware score layer. It records adoption and scores as separate facts, which prevents a benchmark with many mentions from becoming a performance claim.

Suggested citation

Wu, K. (2026). Benchmark Radar: An evidence-first index for tracking AI benchmarks (Technical Report v1.0). DOI: 10.5281/zenodo.22167102

1. System design

Benchmark Radar gives benchmark researchers one place to scan new releases and trace each claim to its source. The collector reads public APIs, feeds, repositories, and scholarly indexes. The pipeline preserves eligible records below the recommendation threshold, so readers can audit the long tail instead of seeing a hand-picked digest.



1.1 Collection and normalization

The configured sources cover scholarly papers, datasets, code, release notes, model cards, and community attention. Required connectors fail the collection run when they break. Optional connectors report health without blocking the other sources. The normalizer maps each item to a shared record shape and retains the source URL, publication time, description, and available artifact identifiers.

Identity rules merge records on exact evidence such as a DOI, arXiv identifier, or repository URL. The system reports unresolved candidates instead of joining items on title similarity alone. This choice leaves some duplicates in the corpus, but it protects distinct benchmarks with similar names.

1.2 Classification and recommendation


A declared taxonomy identifies benchmark, evaluation, dataset, data-quality, and agentic signals. The agentic rule requires an agent term near an evaluation term and excludes survey genres. A scoring rubric ranks records for the daily briefing. In v0.9.0, a score of 40 marks a recommendation; the pipeline still publishes eligible records below 40.

1.3 Two curated measurement layers

Layer	Recorded fact	Guardrail
Model-card adoption	A named vendor document reports a canonical benchmark.	One document adds one mention. A mention carries no score claim.
Score observations	A document prints a numeric result for a named instrument.	Each row stores protocol, source, and reading method. Charts do not connect unrelated points.

2. Cumulative corpus and query surface

The repository stores dated snapshots and replays them into a cumulative evidence graph. Through 29 August 2026, that graph contains 7,540 source observations for 4,537 unique artifacts across 37 snapshots. The normalized benchmark catalog adds 1,173 searchable records. The latest daily snapshot contributes 528 records, of which 186 meet the recommendation threshold. The curated adoption registry contains 94 canonical benchmarks and 36 model cards or release documents. The score layer defines 70 benchmark metrics and records 285 sourced results.



Cumulative observations by source, 37 snapshots through 2026-08-29

2.1 Local retrieval

Installed clients download a checksummed data bundle through explicit **init** and **sync** commands. Search then reads the active local version. The CLI and local HTTP API call the same query service and return the same JSON contract. Version 0.9.0 ranks lexical and token matches; it does not claim semantic retrieval.

Question	Use	Evidence returned
Which benchmarks cover long-horizon agents?	Search the catalog and evidence history.	Matched fields, token coverage, ranking reason, source URLs.
Which benchmarks appear in model cards?	Open the adoption ranking.	Canonical benchmark and the documents that report it.
Did a score rise?	Inspect points that share an instrument and protocol.	Value, date, model, protocol, and cited document.

2.2 Limits

- Public-source coverage changes with API availability, rate limits, and the configured search terms.
- Taxonomy and score rules can miss relevant records or retain false positives. Source links let readers check those decisions.
- Model-card adoption measures vendor reporting choices. Vendors select benchmarks and protocols, so the registry cannot support a neutral model ranking.
- Score comparisons require matching instruments and protocols. Missing protocol details limit the claim to a sourced observation.

3. Reproducibility and access

This report describes Benchmark Radar v0.9.0 at Git commit **98c7de3** and the dated data snapshot **2026-08-29**. The repository publishes source code under the MIT License. Its CI rebuilds normalized external records, classification outputs, and the release bundle before running the test suite.

Artifact	Permanent or canonical location
Technical report	https://doi.org/10.5281/zenodo.22167102
Project and source code	https://github.com/ktwu01/benchmark-radar
Dashboard	https://benchmark-radar.org/
Public JSON corpus	https://benchmark-radar.org/data/radar.json
Machine-readable citation	https://github.com/ktwu01/benchmark-radar/blob/main/CITATION.cff

Data statement

The report statistics come from generated and versioned files in the v0.9.0 repository. **site/data/radar.json** supplies cumulative observation, artifact, source, and snapshot counts after replaying the dated files under **data/snapshots/**. **site/data/benchmark-index.json** supplies the catalog count. **data/model_cards.yml** and **data/benchmark_scores.yml** supply the adoption and score counts.

References

- [1] K. Wu. Benchmark Radar, version 0.9.0. GitHub, 2026. <https://github.com/ktwu01/benchmark-radar>
- [2] K. Wu. Benchmark Radar cumulative corpus schema. 2026. <https://github.com/ktwu01/benchmark-radar/blob/main/docs/cumulative-corpus.schema.json>
- [3] D. S. Katz et al. Recognizing the value of software: a software citation guide. F1000Research 9:1257, 2021. <https://doi.org/10.12688/f1000research.26932.2>
- [4] A. Smith, D. S. Katz, and K. E. Niemeyer. Software citation principles. PeerJ Computer Science 2:e86, 2016. <https://doi.org/10.7717/peerj-cs.86>
- [5] C. E. Jimenez et al. SWE-bench: Can language models resolve real-world GitHub issues? ICLR, 2024. <https://arxiv.org/abs/2310.06770>
- [6] Citation File Format developers. Citation File Format 1.2.0. <https://citation-file-format.github.io/>

Repository: github.com/ktwu01/benchmark-radar
Dashboard: benchmark-radar.org
License: MIT